

# Sistemi Operativi I

Corso di Laurea in Informatica  
2024-2025



**SAPIENZA**  
UNIVERSITÀ DI ROMA

Gabriele Tolomei

Dipartimento di Informatica  
Sapienza Università di Roma

[tolomei@di.uniroma1.it](mailto:tolomei@di.uniroma1.it)

# Contiguous Memory Allocation

- So far, we have assumed each process is allocated into a contiguous space of physical memory

# Contiguous Memory Allocation

- So far, we have assumed each process is allocated into a contiguous space of physical memory
- One simple method is to divide upfront all available memory dedicated to user processes into **equally-sized** segments/partitions

# Contiguous Memory Allocation

- So far, we have assumed each process is allocated into a contiguous space of physical memory
- One simple method is to divide upfront all available memory dedicated to user processes into **equally-sized** segments/partitions
  - Assign each process to a segment

# Contiguous Memory Allocation

- So far, we have assumed each process is allocated into a contiguous space of physical memory
- One simple method is to divide upfront all available memory dedicated to user processes into **equally-sized** segments/partitions
  - Assign each process to a segment
  - Implicitly restricts the grade of multiprogramming (i.e., the number of simultaneous processes) and their size

# Contiguous Memory Allocation


- So far, we have assumed each process is allocated into a contiguous space of physical memory
- One simple method is to divide upfront all available memory dedicated to user processes into **equally-sized** segments/partitions
  - Assign each process to a segment
  - Implicitly restricts the grade of multiprogramming (i.e., the number of simultaneous processes) and their size
  - No longer used!

# Contiguous Memory Allocation

An alternative approach is for the OS to keep track of **free** (unused) memory segments, as processes enter the system, grow, and terminate

# Contiguous Memory Allocation


An alternative approach is for the OS to keep track of **free** (unused) memory segments, as processes enter the system, grow, and terminate





# Contiguous Memory Allocation

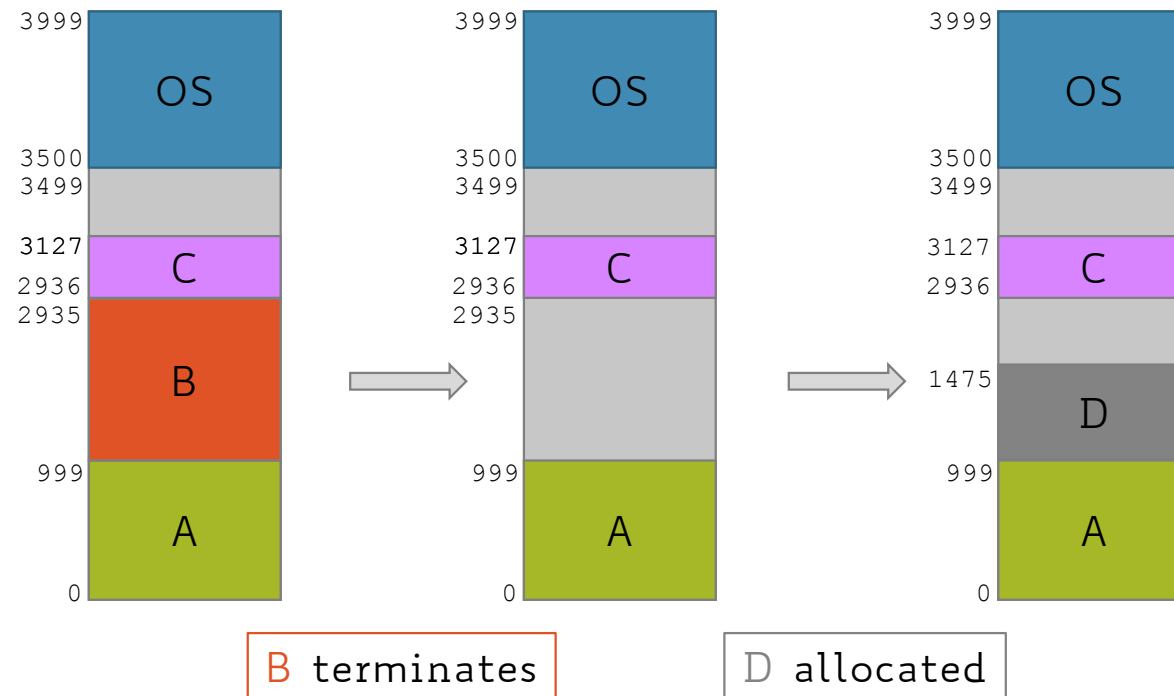
An alternative approach is for the OS to keep track of **free** (unused) memory segments, as processes enter the system, grow, and terminate



B terminates

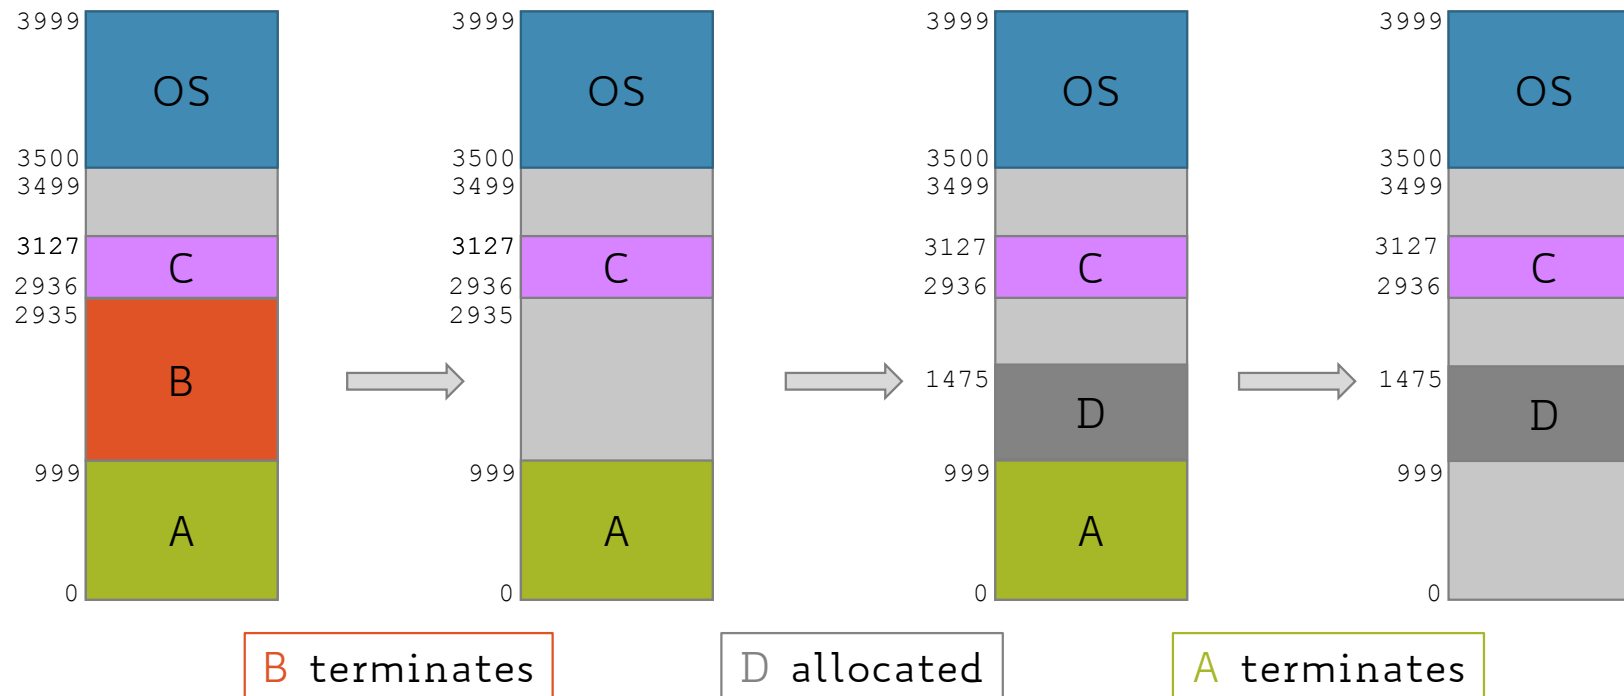
# Contiguous Memory Allocation

An alternative approach is for the OS to keep track of **free** (unused) memory segments, as processes enter the system, grow, and terminate

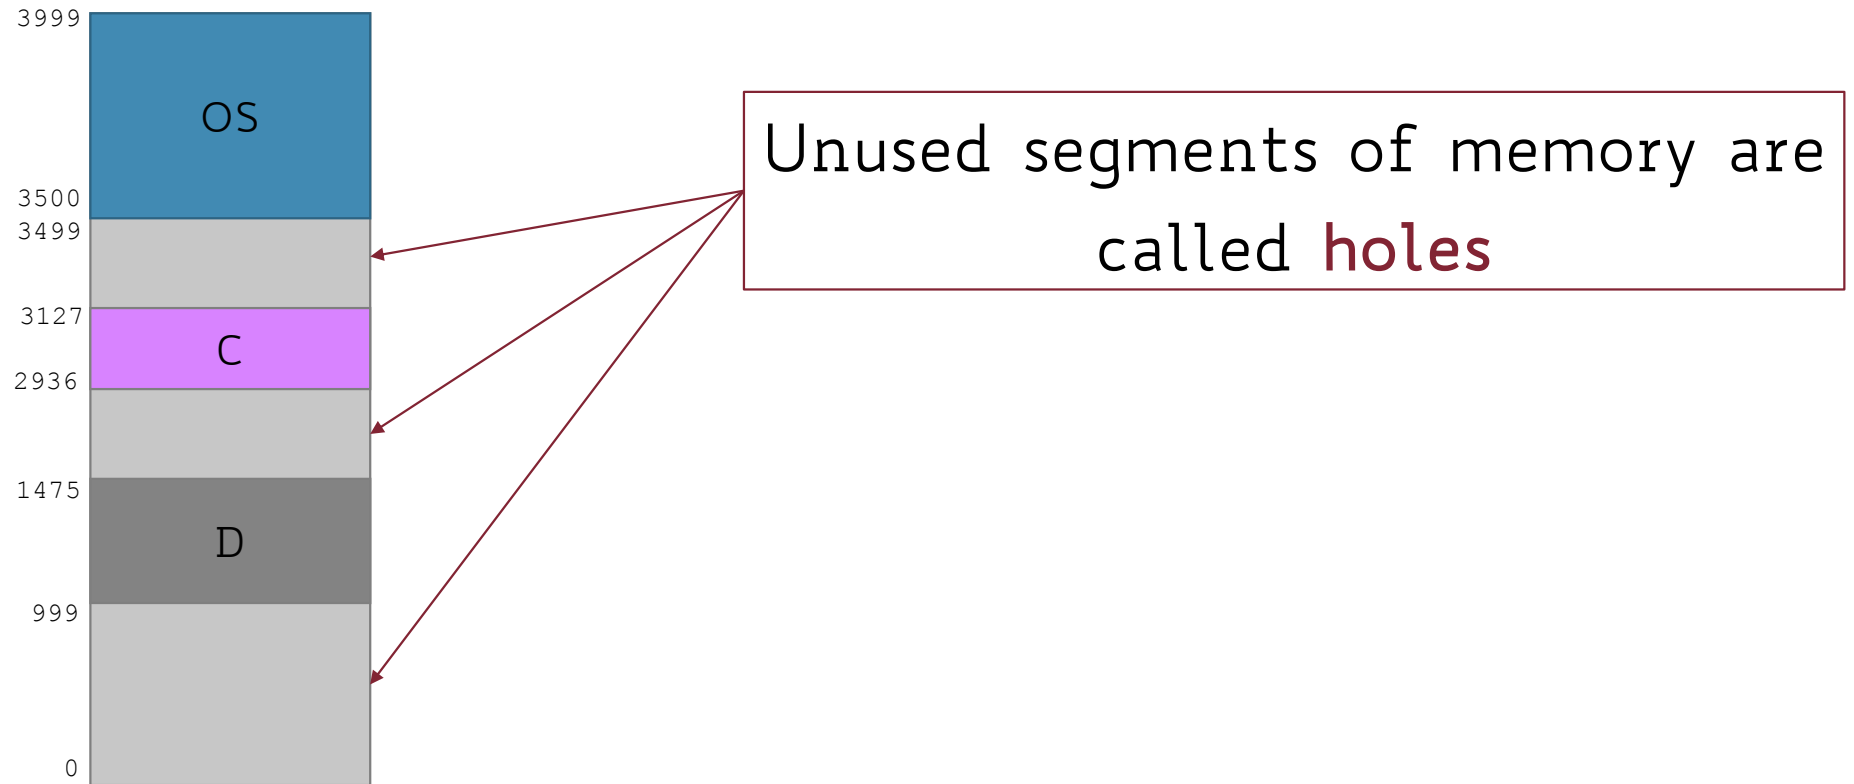


# Contiguous Memory Allocation

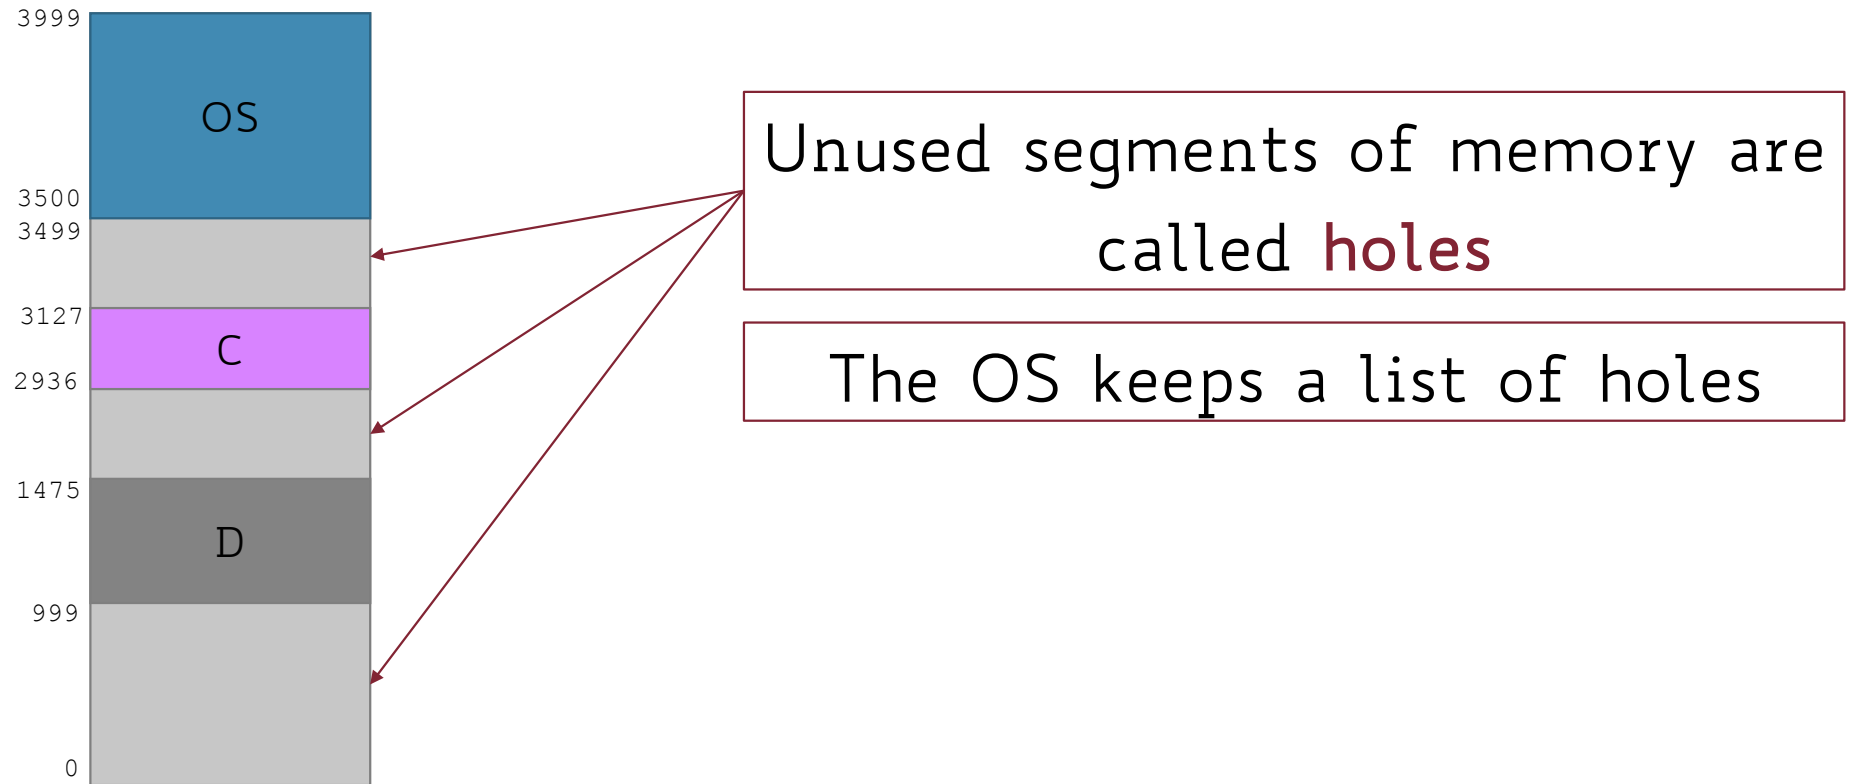
An alternative approach is for the OS to keep track of **free** (unused) memory segments, as processes enter the system, grow, and terminate




# Contiguous Memory Allocation



# Contiguous Memory Allocation



# Contiguous Memory Allocation




Unused segments of memory are called **holes**

The OS keeps a list of holes

Whenever a process has to be loaded, the OS must select a hole of suitable size

# Contiguous Memory Allocation



Unused segments of memory are called **holes**

The OS keeps a list of holes

Whenever a process has to be loaded, the OS must select a hole of suitable size

**How?**

# Memory Allocation Policies: First-Fit

- Linearly scan the list of holes until one is found that is big enough to satisfy the request



# Memory Allocation Policies: First-Fit


- Linearly scan the list of holes until one is found that is big enough to satisfy the request
- Subsequent requests may either start from the beginning of the list or from the end of previous search

# Memory Allocation Policies: First-Fit

- Linearly scan the list of holes until one is found that is big enough to satisfy the request
- Subsequent requests may either start from the beginning of the list or from the end of previous search
- **Complexity:**  $O(n)$ , where  $n$  is the number of holes


# Memory Allocation Policies: First-Fit

Suppose process **X** needs 100B of memory to be loaded, and the list of holes is as follows:



# Memory Allocation Policies: First-Fit


Suppose process X needs 100B of memory to be loaded, and the list of holes is as follows:



Which segment will be picked by the OS using first-fit?

# Memory Allocation Policies: First-Fit


Suppose process X needs 100B of memory to be loaded, and the list of holes is as follows:




Which segment will be picked by the OS using first-fit?

# Memory Allocation Policies: First-Fit

Suppose process **X** needs 100B of memory to be loaded, and the list of holes is as follows:




Which segment will be picked by the OS using first-fit?




# Memory Allocation Policies: First-Fit

Suppose process **X** needs 100B of memory to be loaded, and the list of holes is as follows:




Which segment will be picked by the OS using first-fit?




What if afterwards process **Y** requires 350B?

# Memory Allocation Policies: First-Fit

Suppose process **X** needs 100B of memory to be loaded, and the list of holes is as follows:



Which segment will be picked by the OS using first-fit?



What if afterwards process **Y** requires 350B?

We will not be able to satisfy this request even if theoretically we could



# Memory Allocation Policies: Best-Fit

- Allocate the smallest hole that is big enough to satisfy the request

# Memory Allocation Policies: Best-Fit

- Allocate the smallest hole that is big enough to satisfy the request
- This saves large holes for other process requests that may need them

# Memory Allocation Policies: Best-Fit

- Allocate the smallest hole that is big enough to satisfy the request
- This saves large holes for other process requests that may need them
- However, the resulting unused portions of holes may be too small to be of any use, and will therefore be wasted

# Memory Allocation Policies: Best-Fit

- Allocate the smallest hole that is big enough to satisfy the request
- This saves large holes for other process requests that may need them
- However, the resulting unused portions of holes may be too small to be of any use, and will therefore be wasted
- **Complexity:** still  $O(n)$  but can be  $O(\log n)$  if the list of holes is kept sorted

# Memory Allocation Policies: Best-Fit

- Allocate the smallest hole that is big enough to satisfy the request
- This saves large holes for other process requests that may need them
- However, the resulting unused portions of holes may be too small to be of any use, and will therefore be wasted
- **Complexity:** still  $O(n)$  but can be  $O(\log n)$  if the list of holes is kept sorted

Do you know how which data structure can be used to achieve this?

# Memory Allocation Policies: Best-Fit


- Allocate the smallest hole that is big enough to satisfy the request
- This saves large holes for other process requests that may need them
- However, the resulting unused portions of holes may be too small to be of any use, and will therefore be wasted
- **Complexity:** still  $O(n)$  but can be  $O(\log n)$  if the list of holes is kept sorted

Do you know how which data structure can be used to achieve this?

Binary Search Tree (BST)


# Memory Allocation Policies: Best-Fit

Suppose process **X** needs 100B of memory to be loaded, and the list of holes is as follows:



# Memory Allocation Policies: Best-Fit

Suppose process X needs 100B of memory to be loaded, and the list of holes is as follows:



Which segment will be picked by the OS using best-fit?



# Memory Allocation Policies: Best-Fit

Suppose process X needs 100B of memory to be loaded, and the list of holes is as follows:




Which segment will be picked by the OS using best-fit?

# Memory Allocation Policies: Best-Fit

Suppose process **X** needs 100B of memory to be loaded, and the list of holes is as follows:




Which segment will be picked by the OS using best-fit?



# Memory Allocation Policies: Best-Fit

Suppose process **X** needs 100B of memory to be loaded, and the list of holes is as follows:



Which segment will be picked by the OS using best-fit?




100B only 20B "wasted"

What if afterwards process **Y** requires 350B?

# Memory Allocation Policies: Best-Fit

Suppose process **X** needs 100B of memory to be loaded, and the list of holes is as follows:



Which segment will be picked by the OS using best-fit?



100B only 20B "wasted"

What if afterwards process **Y** requires 350B?

We can now assign it the second available hole segment (400B)

# Memory Allocation Policies: Worst-Fit

- Allocate the largest hole available

# Memory Allocation Policies: Worst-Fit

- Allocate the largest hole available
- Might sound counterintuitive but this increases the likelihood that the remaining portion will be usable for satisfying future requests

# Memory Allocation Policies: Worst-Fit

- Allocate the largest hole available
- Might sound counterintuitive but this increases the likelihood that the remaining portion will be usable for satisfying future requests
- Simulations show that First-Fit and Best-Fit usually work best

# Memory Allocation Policies: Worst-Fit

- Allocate the largest hole available
- Might sound counterintuitive but this increases the likelihood that the remaining portion will be usable for satisfying future requests
- Simulations show that First-Fit and Best-Fit usually work best
- First-Fit is also generally faster than Best-Fit



# Fragmentation

## Problem

Individual holes may be too small to serve a process request but they can be large enough if combined together

# Fragmentation

## Problem

Individual holes may be too small to serve a process request but they can be large enough if combined together



External  
Fragmentation

# Fragmentation

## Problem

Individual holes may be too small to serve a process request but they can be large enough if combined together

```
graph TD; A[Problem] --> B[External Fragmentation]; A --> C[Internal Fragmentation];
```

External  
Fragmentation

Internal  
Fragmentation

# External Fragmentation

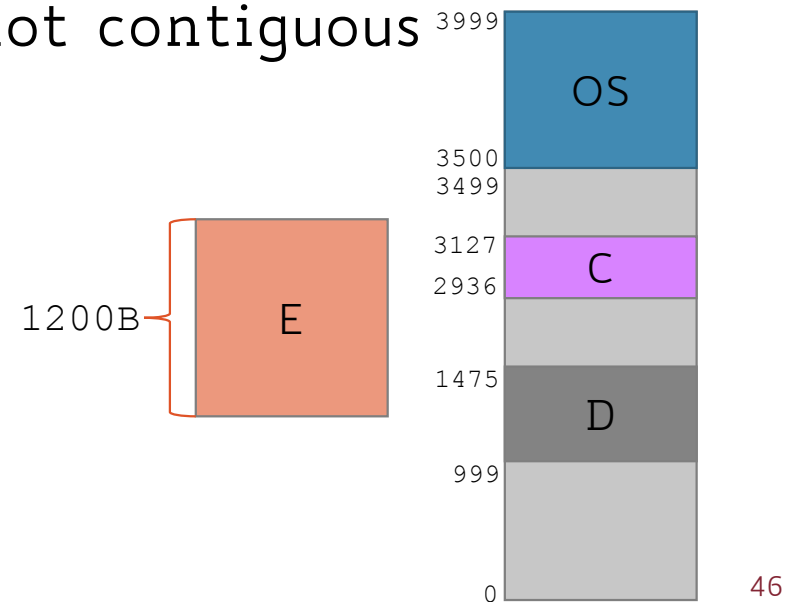
- Frequent loading and unloading processes causes holes to be broken into small (i.e., unusable) chunks

# External Fragmentation

- Frequent loading and unloading processes causes holes to be broken into small (i.e., unusable) chunks
- It happens when there is enough memory to load a process in memory but space is not contiguous

# External Fragmentation

- Frequent loading and unloading processes causes holes to be broken into small (i.e., unusable) chunks
- It happens when there is enough memory to load a process in memory but space is not contiguous

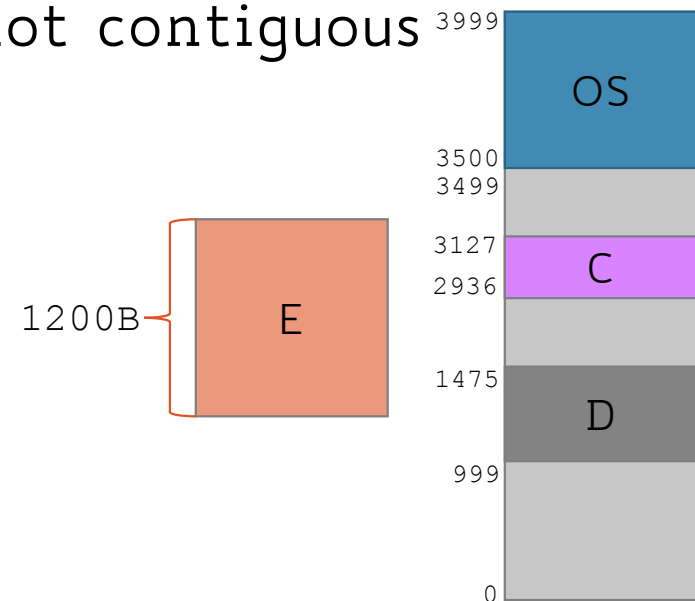


# External Fragmentation

- Frequent loading and unloading processes causes holes to be broken into small (i.e., unusable) chunks
- It happens when there is enough memory to load a process in memory but space is not contiguous

Simulations show that for every 2N allocated blocks, N are lost due to external fragmentation

1/3 of memory space is wasted on average



# External Fragmentation

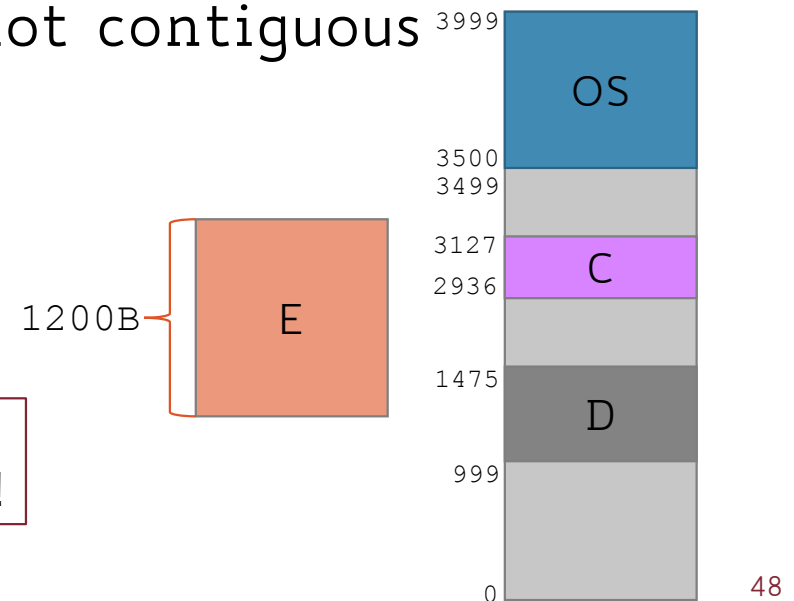
- Frequent loading and unloading processes causes holes to be broken into small (i.e., unusable) chunks
- It happens when there is enough memory to load a process in memory but space is not contiguous

Simulations show that for every 2N allocated blocks, N are lost due to external fragmentation

1/3 of memory space is wasted on average

Goal:

Allocation policy that minimizes wasted space!





# Internal Fragmentation

- It happens when memory internal to a segment is wasted


# Internal Fragmentation

- It happens when memory internal to a segment is wasted
- For example, consider a process whose size is 8,846B and a hole of size 8,848B


# Internal Fragmentation

- It happens when memory internal to a segment is wasted
- For example, consider a process whose size is 8,846B and a hole of size 8,848B
- It may be much more efficient to allocate the process the whole block (and waste 2B) rather than keep track of a tiny 2B hole

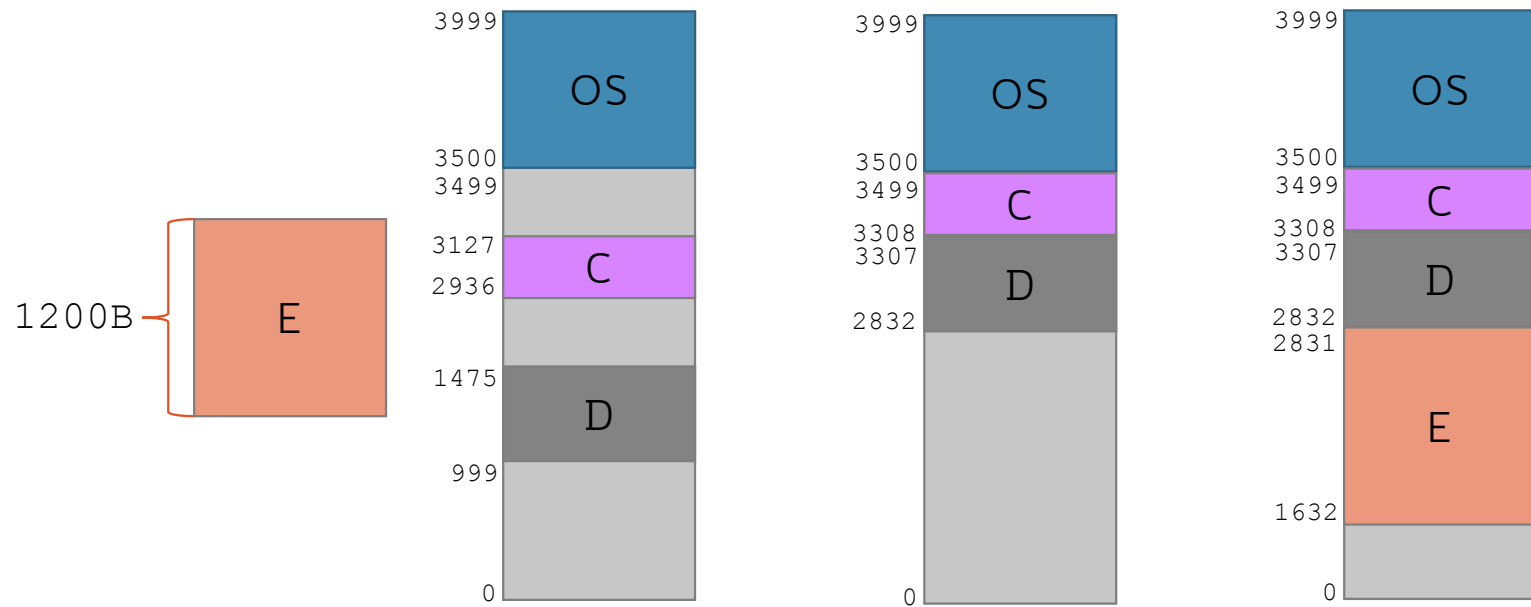
# Solution to Fragmentation: Full Compaction




# Solution to Fragmentation: Full Compaction



# Solution to Fragmentation: Full Compaction




# Solution to Fragmentation: Full Compaction




Only one hole is left but two processes need to be moved (C and D)

# Solution to Fragmentation: Partial Compaction







# Solution to Fragmentation: Partial Compaction



# Solution to Fragmentation: Partial Compaction



# Solution to Fragmentation: Partial Compaction



Still some holes left but only one process is moved (D) rather than two

# Swapping

- So far, we have assumed all processes are entirely loaded in memory (of course, when loaded!)

# Swapping

- So far, we have assumed all processes are entirely loaded in memory (of course, when loaded!)
- **Remember:** A process needs to sit physically in main memory only if the CPU executes its instructions and accesses its data

# Swapping

- So far, we have assumed all processes are entirely loaded in memory (of course, when loaded!)
- **Remember:** A process needs to sit physically in main memory only if the CPU executes its instructions and accesses its data
- If a process blocks (e.g., due to an I/O call) it doesn't need to be in memory while I/O is running

# Swapping

- So far, we have assumed all processes are entirely loaded in memory (of course, when loaded!)
- **Remember:** A process needs to sit physically in main memory only if the CPU executes its instructions and accesses its data
- If a process blocks (e.g., due to an I/O call) it doesn't need to be in memory while I/O is running
- That process can be "swapped out" from memory to disk to make room for other processes

# Swapping

- Once process becomes ready again, the OS must reload it in memory



# Swapping

- Once process becomes ready again, the OS must reload it in memory
- Swap in depends on the address binding used:
  - **compile-** or **load-time**: must be swapped back into the same memory location from which they were swapped out

# Swapping

- Once process becomes ready again, the OS must reload it in memory
- Swap in depends on the address binding used:
  - **compile-** or **load-time**: must be swapped back into the same memory location from which they were swapped out
  - **execution-time**: can be swapped back into any available location (updating base and limit registers)

# Swapping

- Once process becomes ready again, the OS must reload it in memory
- Swap in depends on the address binding used:
  - **compile-** or **load-time**: must be swapped back into the same memory location from which they were swapped out
  - **execution-time**: can be swapped back into any available location (updating base and limit registers)
- Using swapping, fragmentation can be tackled easily
  - Just run compaction before swapping-in a process

# Swapping: Example

- Swapping is a very slow process compared to other operations due to the interaction with hard disk

# Swapping: Example

- Swapping is a very slow process compared to other operations due to the interaction with hard disk
- Example:
  - 10 MB user process
  - disk transfer rate = 40 MB/sec (250 msec just to do the data transfer)


# Swapping: Example

- Swapping is a very slow process compared to other operations due to the interaction with hard disk
- Example:
  - 10 MB user process
  - disk transfer rate = 40 MB/sec (250 msec just to do the data transfer)
- Since swap-in may involve swapping-out another process, the overall time required will be ~500 msec

# Swapping: Example

- Swapping is a very slow process compared to other operations due to the interaction with hard disk
- Example:
  - 10 MB user process
  - disk transfer rate = 40 MB/sec (250 msec just to do the data transfer)
- Since swap-in may involve swapping-out another process, the overall time required will be ~500 msec
- Time slice is usually way smaller than that!

# Swapping



Most modern OSs no longer use swapping, because it is too slow and there are faster alternatives available (e.g., **paging**)



# Problems Seen So Far

- Contiguous allocation
  - Hard to grow or shrink process memory

# Problems Seen So Far

- Contiguous allocation
  - Hard to grow or shrink process memory
- Fragmentation
  - Frequent compaction needed

# Problems Seen So Far

- Contiguous allocation
  - Hard to grow or shrink process memory
- Fragmentation
  - Frequent compaction needed
- Process entirely loaded
  - Swapping helps but it may be too inefficient

# Paging

- A memory management scheme that addresses the problems above

# Paging

- A memory management scheme that addresses the problems above
- The logical address space of a process is still **contiguous** but it is divided into **fixed-size blocks**, called **pages**

# Paging

- A memory management scheme that addresses the problems above
- The logical address space of a process is still **contiguous** but it is divided into **fixed-size blocks**, called **pages**
- Contiguous allocation is no longer required as logical pages can be mapped to **non-contiguous** physical **frames**

# Paging

- A memory management scheme that addresses the problems above
- The logical address space of a process is still **contiguous** but it is divided into **fixed-size blocks**, called **pages**
- Contiguous allocation is no longer required as logical pages can be mapped to **non-contiguous** physical **frames**
- External fragmentation is eliminated because pages have fixed size
  - Internal fragmentation may still occur though

# Paging

- A memory management scheme that addresses the problems above
- The logical address space of a process is still **contiguous** but it is divided into **fixed-size blocks**, called **pages**
- Contiguous allocation is no longer required as logical pages can be mapped to **non-contiguous** physical **frames**
- External fragmentation is eliminated because pages have fixed size
  - Internal fragmentation may still occur though

## 90/10 Rule

Processes spend **90%** of their time accessing only **10%** of their allocated memory space

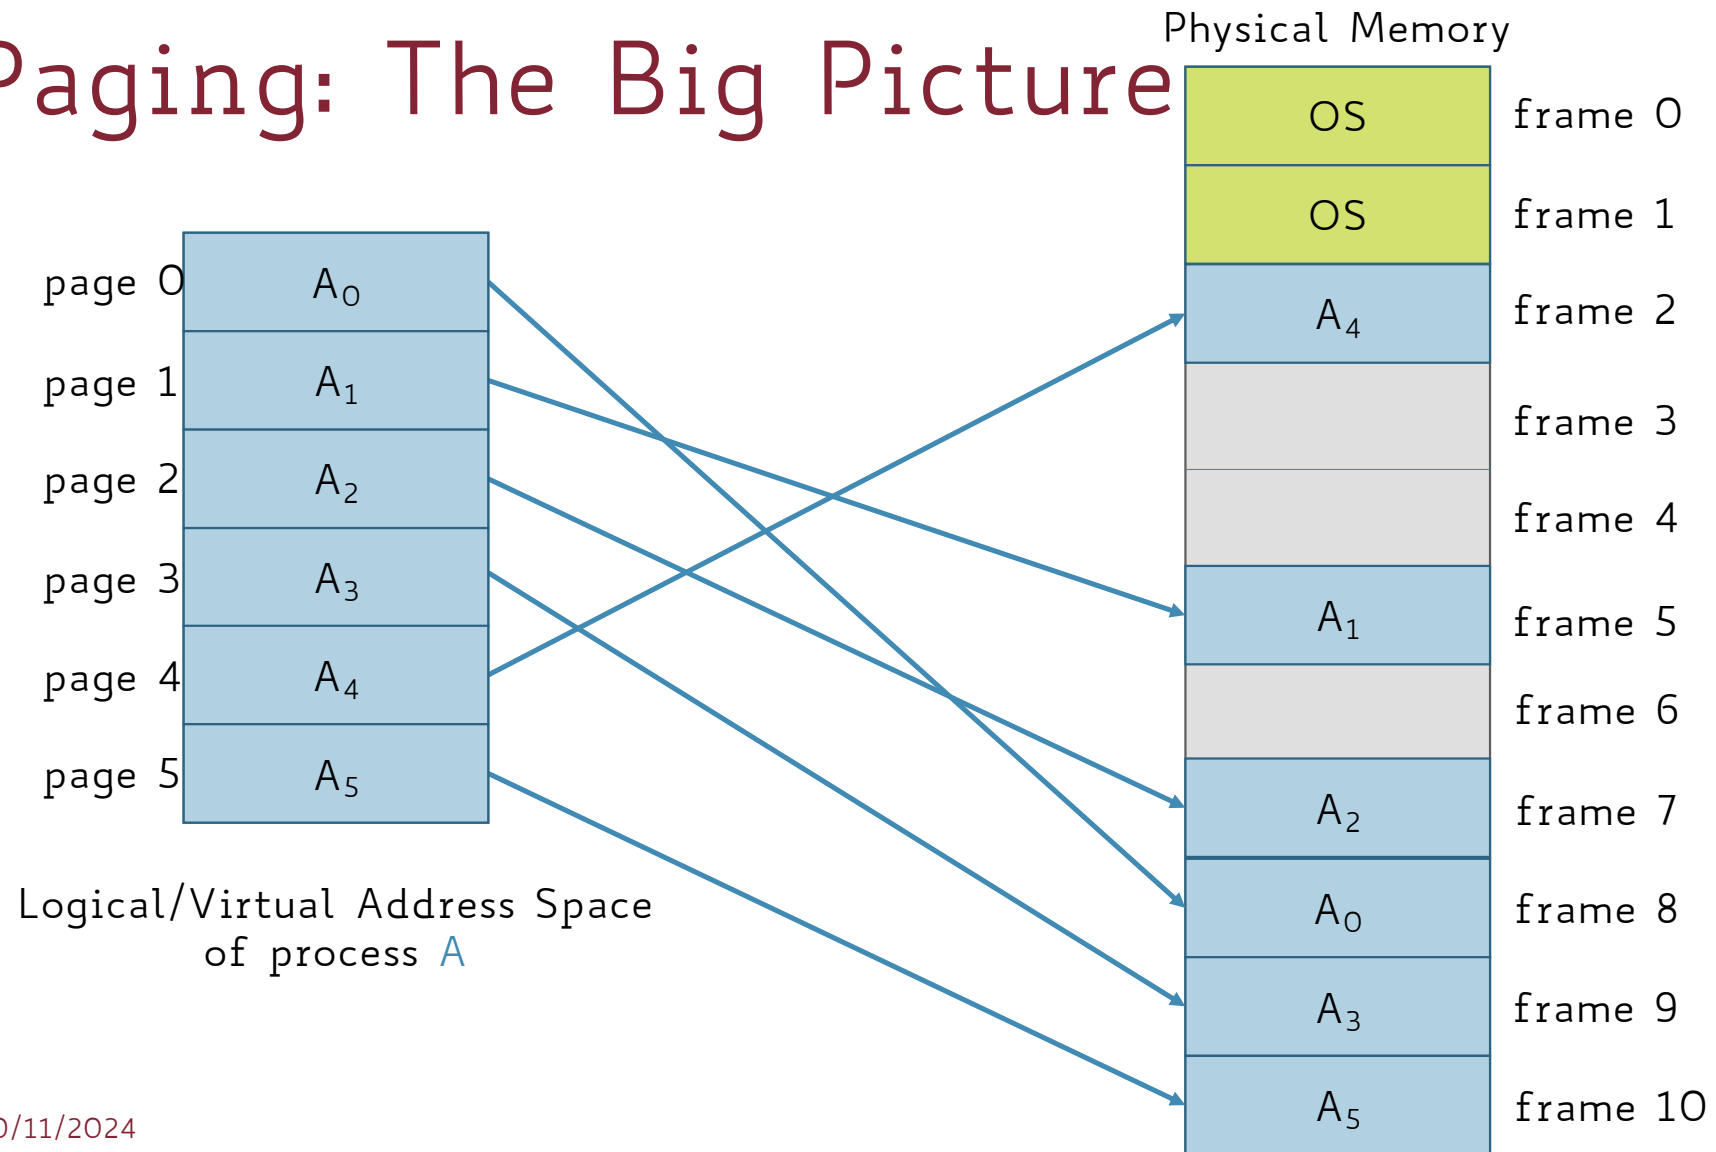


# Paging: The Big Picture

page 0	$A_0$
page 1	$A_1$
page 2	$A_2$
page 3	$A_3$
page 4	$A_4$
page 5	$A_5$

Logical/Virtual Address Space  
of process  $A$

# Paging: The Big Picture



# Basic OS Responsibilities for Paging

- The OS has 2 main responsibilities:
  - mapping between logical pages and physical frames
  - translating logical addresses to physical addresses

# Basic OS Responsibilities for Paging

- The OS has 2 main responsibilities:
  - mapping between logical pages and physical frames
  - translating logical addresses to physical addresses
- All of this must be done efficiently!
  - Remember, memory addresses are referenced all the time

# Basic OS Responsibilities for Paging

- The OS has 2 main responsibilities:
  - mapping between logical pages and physical frames
  - translating logical addresses to physical addresses
- All of this must be done efficiently!
  - Remember, memory addresses are referenced all the time
- OS needs dedicated support for doing it → **Page Table**

# Page Table: Mapping Pages to Frames

0	$A_0$
1	$A_1$
2	$A_2$
3	$A_3$
4	$A_4$
5	$A_5$

OS	0
OS	1
$A_4$	2
	3
	4
$A_1$	5
	6
$A_2$	7
$A_0$	8
$A_3$	9
$A_5$	10

# Page Table: Mapping Pages to Frames

Lookup table to retrieve what frame a page is stored in

0	A <sub>0</sub>
1	A <sub>1</sub>
2	A <sub>2</sub>
3	A <sub>3</sub>
4	A <sub>4</sub>
5	A <sub>5</sub>

Page	Frame
0	8
1	5
2	7
3	9
4	2
5	10

OS	0
OS	1
A <sub>4</sub>	2
	3
	4
A <sub>1</sub>	5
	6
A <sub>2</sub>	7
A <sub>0</sub>	8
A <sub>3</sub>	9
A <sub>5</sub>	10

# Page Table: Mapping Pages to Frames

Lookup table to retrieve what frame a page is stored in

0	A <sub>0</sub>
1	A <sub>1</sub>
2	A <sub>2</sub>
3	A <sub>3</sub>
4	A <sub>4</sub>
5	A <sub>5</sub>

Page	Frame
0	8
1	5
2	7
3	9
4	2
5	10

OS	0
OS	1
A <sub>4</sub>	2
	3
	4
A <sub>1</sub>	5
	6
A <sub>2</sub>	7
A <sub>0</sub>	8
A <sub>3</sub>	9
A <sub>5</sub>	10

We have assumed **all** pages of a process are mapped to physical frames, but this is not always the case



# Summary

- Contiguous memory allocation may cause fragmentation
  - External vs. Internal Fragmentation

# Summary

- Contiguous memory allocation may cause fragmentation
  - External vs. Internal Fragmentation
- Existing countermeasures (compaction) exist but they are costly

# Summary

- Contiguous memory allocation may cause fragmentation
  - External vs. Internal Fragmentation
- Existing countermeasures (compaction) exist but they are costly
- We may want to relax the constraint on having an entire process loaded in main memory

# Summary

- Contiguous memory allocation may cause fragmentation
  - External vs. Internal Fragmentation
- Existing countermeasures (compaction) exist but they are costly
- We may want to relax the constraint on having an entire process loaded in main memory
- Paging solves all these issues!