

# Sistemi Operativi I

Corso di Laurea in Informatica  
2023-2024



SAPIENZA  
UNIVERSITÀ DI ROMA

Gabriele Tolomei

Dipartimento di Informatica

Sapienza Università di Roma

[tolomei@di.uniroma1.it](mailto:tolomei@di.uniroma1.it)

# Paging + Segmentation

- Paging (OS' view of memory)
  - Divide memory into fixed-size pages and map them to physical frames
- Segmentation (compiler's view of memory)
  - Divide process into logical segments (e.g., code, data, stack, heap)
- Combine paging with segmentation
  - Segmented Paging

# Paging + Segmentation

- Paging (OS' view of memory)
  - Divide memory into fixed-size pages and map them to physical frames
- Segmentation (compiler's view of memory)
  - Divide process into logical segments (e.g., code, data, stack, heap)
- Combine paging with segmentation
  - Segmented Paging

So far, the entire virtual address space of a process was assumed to fit and be all in memory

# Virtual Memory

- In practice, most real processes do not need all their pages loaded in memory, or at least not all at once, e.g.,:

# Virtual Memory

- In practice, most real processes do not need all their pages loaded in memory, or at least not all at once, e.g.,:
  - Error handling code is not needed unless that specific error occurs, some of which are quite rare

# Virtual Memory

- In practice, most real processes do not need all their pages loaded in memory, or at least not all at once, e.g.,:
  - Error handling code is not needed unless that specific error occurs, some of which are quite rare
  - Arrays are often over-sized for worst-case scenarios, and only a small fraction of the arrays is actually used in practice

# Virtual Memory

- In practice, most real processes do not need all their pages loaded in memory, or at least not all at once, e.g.,:
  - Error handling code is not needed unless that specific error occurs, some of which are quite rare
  - Arrays are often over-sized for worst-case scenarios, and only a small fraction of the arrays is actually used in practice
  - Some features of certain programs are rarely used

# Virtual Memory

- In practice, most real processes do not need all their pages loaded in memory, or at least not all at once, e.g.,:
  - Error handling code is not needed unless that specific error occurs, some of which are quite rare
  - Arrays are often over-sized for worst-case scenarios, and only a small fraction of the arrays is actually used in practice
  - Some features of certain programs are rarely used

**Virtual Memory** uses backing storage (i.e., disk) to store unused pages and give the illusion of infinite virtual address space



# Virtual Memory: Benefits

- The ability to load only the portions of processes that are actually needed (and only when needed) from disk has several benefits:

# Virtual Memory: Benefits

- The ability to load only the portions of processes that are actually needed (and only when needed) from disk has several benefits:
  - Programs could be written for a much larger address space than physically exists on the computer


# Virtual Memory: Benefits

- The ability to load only the portions of processes that are actually needed (and only when needed) from disk has several benefits:
  - Programs could be written for a much larger address space than physically exists on the computer
  - More memory is left for other programs, improving CPU utilization

# Virtual Memory: Benefits


- The ability to load only the portions of processes that are actually needed (and only when needed) from disk has several benefits:
  - Programs could be written for a much larger address space than physically exists on the computer
  - More memory is left for other programs, improving CPU utilization
  - Less I/O is needed for swapping processes in and out of memory, speeding things up

# Virtual Memory: The Big Picture




At any given time,  
each **page** can be:

# Virtual Memory: The Big Picture



At any given time,  
each **page** can be:  
in **memory** (physical frame)


# Virtual Memory: The Big Picture



At any given time,  
each **page** can be:  
in **memory** (physical frame)  
on **backing store** (disk)

# Virtual Memory: The Big Picture


virtual memory can be much larger than physical memory



At any given time,  
each **page** can be:  
in **memory** (physical frame)  
on **backing store** (disk)




# The Sparseness of Virtual Address Space



Typically, virtual address space is highly sparse

# The Sparseness of Virtual Address Space



Typically, virtual address space is highly sparse

A lot of virtual memory addresses remain unreferenced

# Virtual Memory: Idea

- Use the memory as a cache for the disk

# Virtual Memory: Idea

- Use the memory as a cache for the disk
- The page table must also indicate if the page is on disk or in memory (just using a single invalid bit)

# Virtual Memory: Idea

- Use the memory as a cache for the disk
- The page table must also indicate if the page is on disk or in memory (just using a single invalid bit)
- Once the page is loaded from disk to memory, the OS updates the corresponding entry of the page table along with the valid bit

# Virtual Memory: Idea

- Remember: access to disk is extremely slower than access to memory

# Virtual Memory: Idea

- Remember: access to disk is extremely slower than access to memory
- Therefore, memory accesses must reference pages that are in memory **with high probability**

# Virtual Memory: The Locality Principle

- The 90÷10 rule claims that on a particular time frame, most of the memory references made by a process is around a small "area"



# Virtual Memory: The Locality Principle

- The 90÷10 rule claims that on a particular time frame, most of the memory references made by a process is around a small "area"
- We call this area as the **working set** of the process

# Virtual Memory: The Locality Principle

- The 90÷10 rule claims that on a particular time frame, most of the memory references made by a process is around a small "area"
- We call this area as the **working set** of the process
- Since the working set is fairly small compared to the whole virtual address space, it will likely fit in memory


# Virtual Memory: The Locality Principle

- Of course, during the lifetime of a process its working set may change (i.e., a process may eventually refer *all* of its virtual address space)

# Virtual Memory: The Locality Principle


- Of course, during the lifetime of a process its working set may change (i.e., a process may eventually refer *all* of its virtual address space)
- But in a reasonably small time frame, the working set stays "the same"

# Virtual Memory: Basic Concepts



At each logical memory reference, a page table lookup is performed as usual


# Virtual Memory: Basic Concepts



At each logical memory reference, a page table lookup is performed as usual

In the page table, the valid-invalid bit is checked for the corresponding entry

# Virtual Memory: Basic Concepts




At each logical memory reference, a page table lookup is performed as usual

In the page table, the valid-invalid bit is checked for the corresponding entry

If the bit is set to 1 it means the page entry is valid (i.e., the requested page is in memory)

# Virtual Memory: Basic Concepts



At each logical memory reference, a page table lookup is performed as usual

In the page table, the valid-invalid bit is checked for the corresponding entry


If the bit is set to 1 it means the page entry is valid (i.e., the requested page is in memory)

Otherwise, a **page fault trap** occurs, and the page has to be loaded (i.e., fetched) from disk



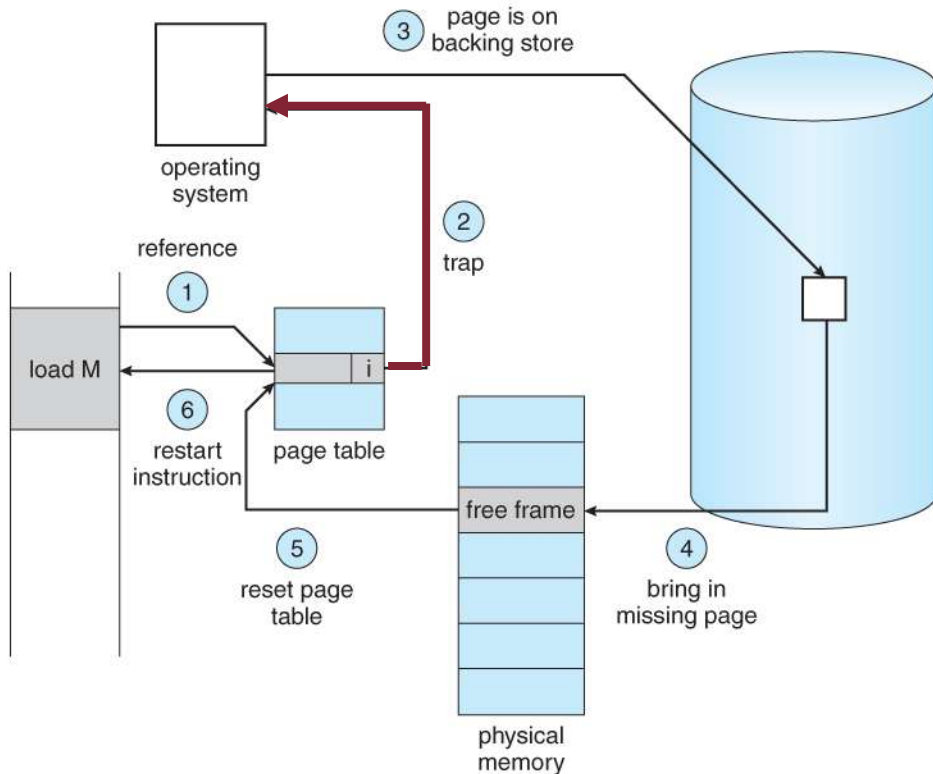
# Page Fault Handling

1. The memory address is first checked, to see if it is legitimate

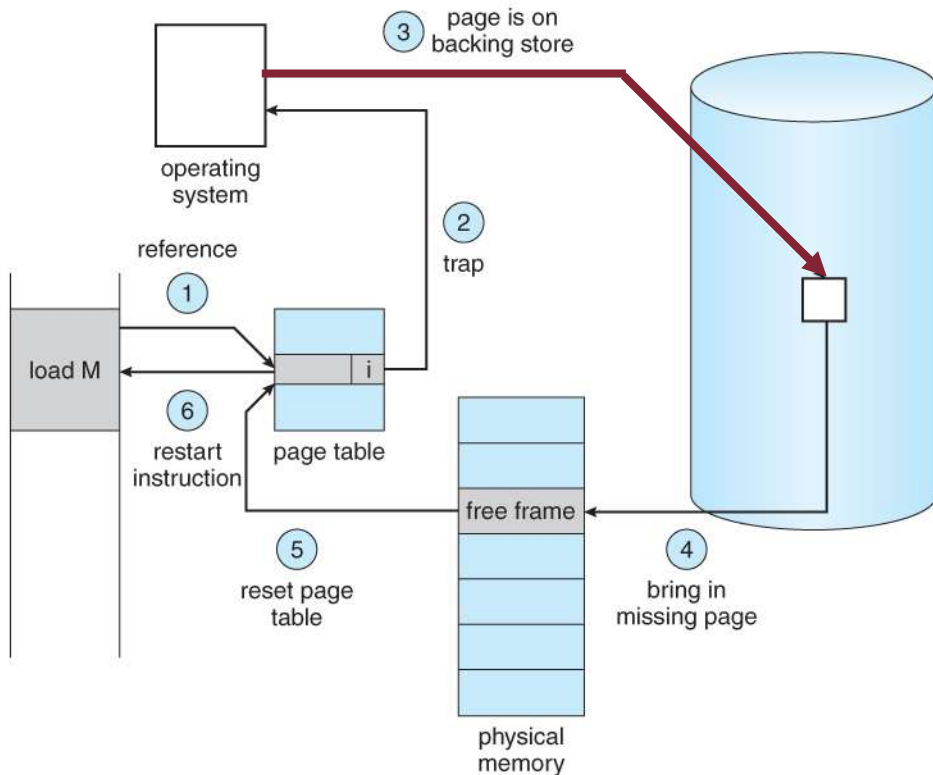


# Page Fault Handling

1. The memory address is first checked, to see if it is legitimate
2. If the address is legitimate the page must be fetched from disk

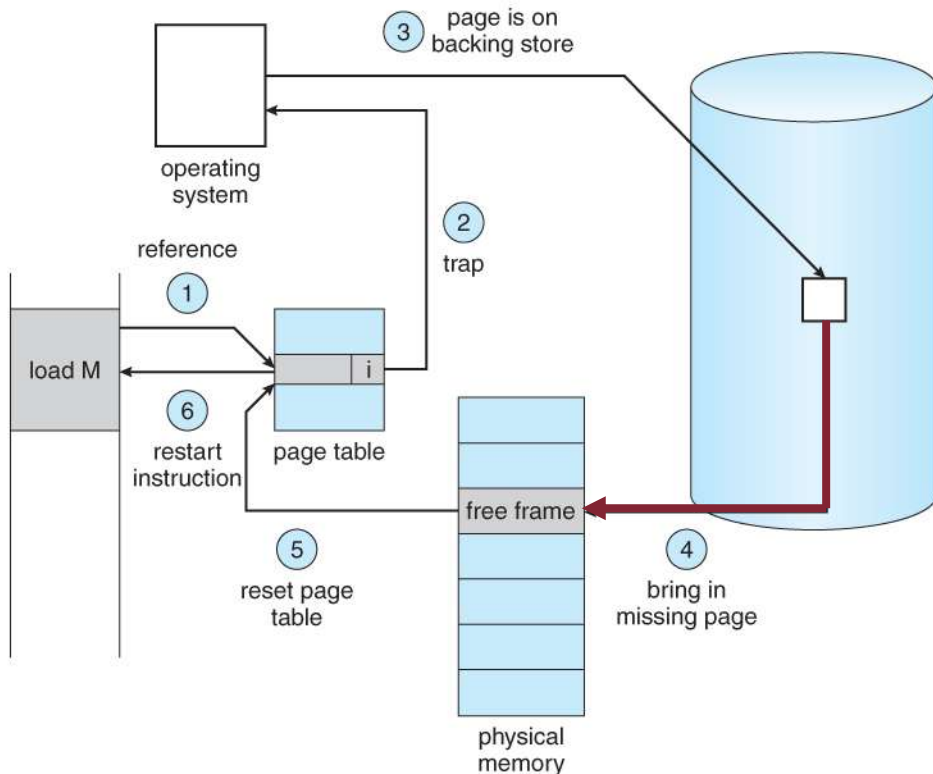


# Page Fault Handling



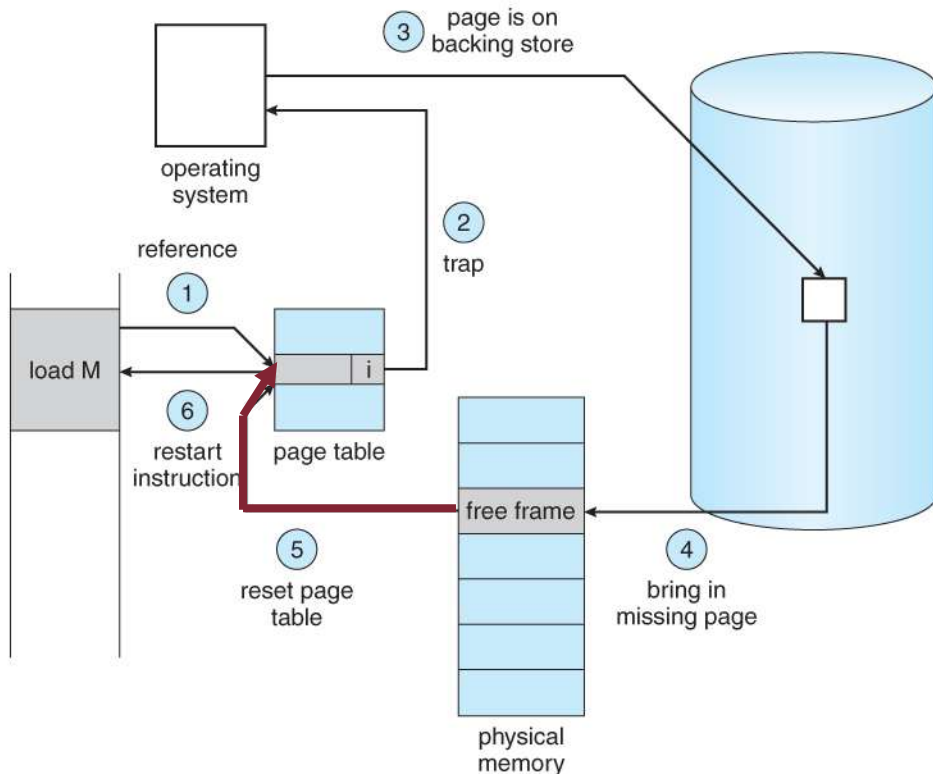
1. The memory address is first checked, to see if it is legitimate
2. If the address is legitimate the page must be fetched from disk
3. A free frame is located, possibly from a free-frame list (the OS might need to pick a frame to unload if all memory is full)

# Page Fault Handling




1. The memory address is first checked, to see if it is legitimate
2. If the address is legitimate the page must be fetched from disk
3. A free frame is located, possibly from a free-frame list (the OS might need to pick a frame to unload if all memory is full)
4. A disk operation is scheduled to bring in the necessary page from disk (this will block the process on an I/O wait, allowing some other process to run)

# Page Fault Handling



1. The memory address is first checked, to see if it is legitimate
2. If the address is legitimate the page must be fetched from disk
3. A free frame is located, possibly from a free-frame list (the OS might need to pick a frame to unload if all memory is full)
4. A disk operation is scheduled to bring in the necessary page from disk (this will block the process on an I/O wait, allowing some other process to run)
5. When the I/O operation is complete, the process's page table is updated with the new frame number, and the bit is set to valid

# Page Fault Handling



1. The memory address is first checked, to see if it is legitimate
2. If the address is legitimate the page must be fetched from disk
3. A free frame is located, possibly from a free-frame list (the OS might need to pick a frame to unload if all memory is full)
4. A disk operation is scheduled to bring in the necessary page from disk (this will block the process on an I/O wait, allowing some other process to run)
5. When the I/O operation is complete, the process's page table is updated with the new frame number, and the bit is set to valid
6. The current process gets interrupted and the instruction that caused the page fault must be restarted from the beginning

# Page Fault Handling: TLB Hit

- The TLB also uses the valid bit to indicate the fact that the page is in main memory

# Page Fault Handling: TLB Hit

- The TLB also uses the valid bit to indicate the fact that the page is in main memory
- If we get a TLB hit but the frame is not actually in main memory, we have to go fetch the page from disk anyway!



# Page Fault Handling: TLB Hit

- The TLB also uses the valid bit to indicate the fact that the page is in main memory
- If we get a TLB hit but the frame is not actually in main memory, we have to go fetch the page from disk anyway!
- TLB hit means the requested page entry is in the cache **and** the referenced frame is also in memory

# Page Fault Handling: TLB Miss

- If the requested page is not in the cache (TLB miss) but it is in memory:

# Page Fault Handling: TLB Miss

- If the requested page is not in the cache (TLB miss) but it is in memory:
  - The OS picks a TLB entry to replace and fills it with the new entry

# Page Fault Handling: TLB Miss

- If the requested page is not in the cache (TLB miss) but it is in memory:
  - The OS picks a TLB entry to replace and fills it with the new entry

# Page Fault Handling: TLB Miss

- If the requested page is not in the cache (TLB miss) and it is not even in memory (i.e., it is sitting on disk):
  - The OS picks a TLB entry to replace and fills it with the new entry as follows
    - invalidates the TLB entry
    - performs page fault trap operations
    - updates the TLB entry
    - restarts the faulting instruction

# Page Fault Handling: Faulty Address

- How does the OS figure out which page generated the fault?

# Page Fault Handling: Faulty Address

- How does the OS figure out which page generated the fault?
- Architecture-dependent:
  - x86: hardware saves the virtual address that caused the fault (CR2 register)
  - On some platforms, OS gets only address of faulting instruction, must simulate the instruction and try every address to find the one that generated the fault

# Page Fault Handling: Transparency

- Transparently restarting process execution after a page fault is tricky, since the fault may have occurred in the middle of an instruction



# Page Fault Handling: Transparency

- Transparently restarting process execution after a page fault is tricky, since the fault may have occurred in the middle of an instruction
- To restart (from scratch) a faulty instruction the OS needs hardware support for saving:
  - The faulting instruction
  - The CPU state

# Page Fault Handling: Transparency

- idempotent vs. non-idempotent instructions

# Page Fault Handling: Transparency

- `idempotent` vs. `non-idempotent` instructions
- `idempotent` → just restart the faulting instruction  
(hardware saves instruction address during page fault)

# Page Fault Handling: Transparency

- **idempotent** vs. **non-idempotent** instructions
- **idempotent** → just restart the faulting instruction  
(hardware saves instruction address during page fault)
- **non-idempotent** → much more difficult to restart
  - `MOV [%R1], +(%R2)` → increment the value of R2 and store it to memory address in R1
  - What if memory address [%R1] causes the page fault?
  - Cannot naively redo the instruction from scratch, otherwise R2 gets incremented twice

# Page Fault Handling: Transparency

- Even harder when using instructions that are not easily undoable
  - E.g., instructions that are used to move a block of memory at once
  - The block may span multiple pages: some of them can be in memory while some others not
  - Pages that are in memory can be changed meanwhile a page fault occurs

# Page Fault Handling: Transparency

- Even harder when using instructions that are not easily undoable
  - E.g., instructions that are used to move a block of memory at once
  - The block may span multiple pages: some of them can be in memory while some others not
  - Pages that are in memory can be changed meanwhile a page fault occurs

How to unwind those complicated side-effects?

# Page Fault Handling: Transparency

- Even harder when using instructions that are not easily undoable
  - E.g., instructions that are used to move a block of memory at once
  - The block may span multiple pages: some of them can be in memory while some others not
  - Pages that are in memory can be changed meanwhile a page fault occurs

Ensure all the addresses within the block to be moved are in memory before executing the instruction

# Virtual Memory: Performance

- Theoretically, a page fault may occur at each process instruction
  - A process may reference addresses belonging to different page at each step



# Virtual Memory: Performance

- Theoretically, a page fault may occur at each process instruction
  - A process may reference addresses belonging to different page at each step
- Luckily, processes usually exhibit so-called **locality of reference**

# Virtual Memory: Performance

- Theoretically, a page fault may occur at each process instruction
  - A process may reference addresses belonging to different page at each step
- Luckily, processes usually exhibit so-called **locality of reference**
  - **temporal** → if a process accesses an item in memory, it will tend to reference the same item again soon

# Virtual Memory: Performance

- Theoretically, a page fault may occur at each process instruction
  - A process may reference addresses belonging to different page at each step
- Luckily, processes usually exhibit so-called **locality of reference**
  - **temporal** → if a process accesses an item in memory, it will tend to reference the same item again soon
  - **spatial** → if a process accesses an item in memory, it will tend to reference a close item again soon

# Virtual Memory: Performance

$t_{MA}$  = physical memory access time

$t_{FAULT}$  = time to handle a page fault

$p \in [0, 1]$  = probability of page fault

$t_{ACCESS}$  = effective time for each memory reference

$$t_{ACCESS} = (1 - p) * t_{MA} + p * t_{FAULT}$$

Let's assume:  $t_{MA} = 100$  nsec and  $t_{FAULT} = 20$  msec = 20,000,000 nsec

$$t_{ACCESS} = (1 - p) * 100 + p * 20,000,000$$

# Virtual Memory: Performance

$t_{MA}$  = physical memory access time

$t_{FAULT}$  = time to handle a page fault

$p \in [0, 1]$  = probability of page fault

$t_{ACCESS}$  = effective time for each memory reference

$$t_{ACCESS} = (1 - p) * t_{MA} + p * t_{FAULT}$$

Let's assume:  $t_{MA} = 100$  nsec and  $t_{FAULT} = 20$  msec = 20,000,000 nsec

$$t_{ACCESS} = (1 - p) * 100 + p * 20,000,000$$

This heavily depends on  $p$ !

# Virtual Memory: Performance Example

$$t_{ACCESS} = (1 - p) * 100 + p * 20,000,000$$

What if only 1 every 1,000 memory references causes a page fault  
(i.e.,  $p = 0.001$ )

# Virtual Memory: Performance Example

$$t_{ACCESS} = (1 - p) * 100 + p * 20,000,000$$

What if only 1 every 1,000 memory references causes a page fault  
(i.e.,  $p = 0.001$ )

The access time increases from just 100 nsec up to ~20.1 microsec

200 times slowdown factor

# Virtual Memory: Performance Example

$$t_{ACCESS} = (1 - p) * 100 + p * 20,000,000$$

What if we want the time access to be at most 10% slower than basic memory access?



# Virtual Memory: Performance Example

$$t_{ACCESS} = (1 - p) * 100 + p * 20,000,000$$

What if we want the time access to be at most 10% slower than basic memory access?

We have to solve for  $p$  the following equation:

$$1.1 * 100 = (1 - p) * 100 + p * 20,000,000$$

# Virtual Memory: Performance Example

$$t_{ACCESS} = (1 - p) * 100 + p * 20,000,000$$

What if we want the time access to be at most 10% slower than basic memory access?

We have to solve for  $p$  the following equation:

$$1.1 * 100 = (1 - p) * 100 + p * 20,000,000$$

$$1.1 * 100 = 100 - 100p + 20,000,000p = 19,999,900p = 110 - 100 =$$

$$p = \frac{10}{19,999,900} = \frac{1}{1,999,990} \approx 0,0000005 = 5 * 10^{-7}$$

To achieve that goal, we can tolerate at most 1 page fault every about 2 million accesses!

# Virtual Memory: Performance Example

More generally, given  $t_{MA}$ ,  $t_{FAULT}$ , and a threshold  $\epsilon > 0$  if we want to find  $p$  s.t.:

$$t_{ACCESS} = (1 + \epsilon) * t_{MA}$$

We substitute  $t_{ACCESS}$  and solve for  $p$  the resulting equation:

$$\begin{aligned}(1 - p) * t_{MA} + p * t_{FAULT} &= (1 + \epsilon) * t_{MA} = \\ t_{MA} - p * t_{MA} + p * t_{FAULT} &= t_{MA} + \epsilon * t_{MA} \\ p(t_{FAULT} - t_{MA}) &= \epsilon * t_{MA} =\end{aligned}$$

$$p = \frac{\epsilon * t_{MA}}{t_{FAULT} - t_{MA}}$$

# Virtual Memory: Considerations

- So far, we have described how the OS (with the support of HW) manages page faults

# Virtual Memory: Considerations

- So far, we have described how the OS (with the support of HW) manages page faults
- Still, the OS has to answer 2 fundamental questions:
  - When to load process' pages into main memory (**page fetching**)
  - Which page to remove from memory if this gets filled (**page replacement**)

# Page Fetching Goals

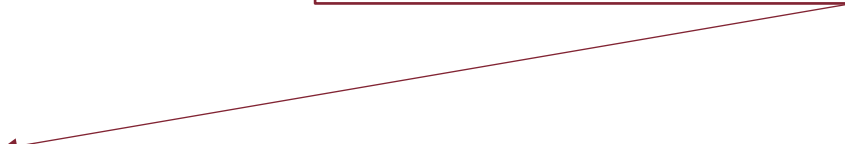
- The overall goal is still to make physical memory look larger than it is
- Exploiting the locality reference of programs
- Keep in memory only those pages that is being used
- Keep on disk those pages that are unused
- Ideally, producing a memory system with the performance of main memory and the cost/capacity of disk!

# Page Fetching Strategies

3 page fetching strategies

# Page Fetching Strategies

3 page fetching strategies



## Startup

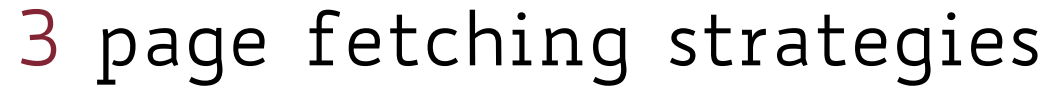
This is a special case where all the pages of the process are loaded at once

virtual address space cannot be larger than physical memory



# Page Fetching Strategies

3 page fetching strategies



```
graph TD; A[3 page fetching strategies] --> B[Startup]; A --> C[Overlays];
```

## Startup

This is a special case where all the pages of the process are loaded at once

virtual address space cannot be larger than physical memory

## Overlays

Let the programmer say when pages are loaded/removed

virtual address space can be larger than physical memory but hard and error-prone

# Page Fetching Strategies

3 page fetching strategies

```
graph TD; A[3 page fetching strategies] --> B[Startup]; A --> C[Overlays]; A --> D[Demand];
```

## Startup

This is a special case where all the pages of the process are loaded at once

virtual address space cannot be larger than physical memory

## Overlays

Let the programmer say when pages are loaded/removed

virtual address space can be larger than physical memory but hard and error-prone

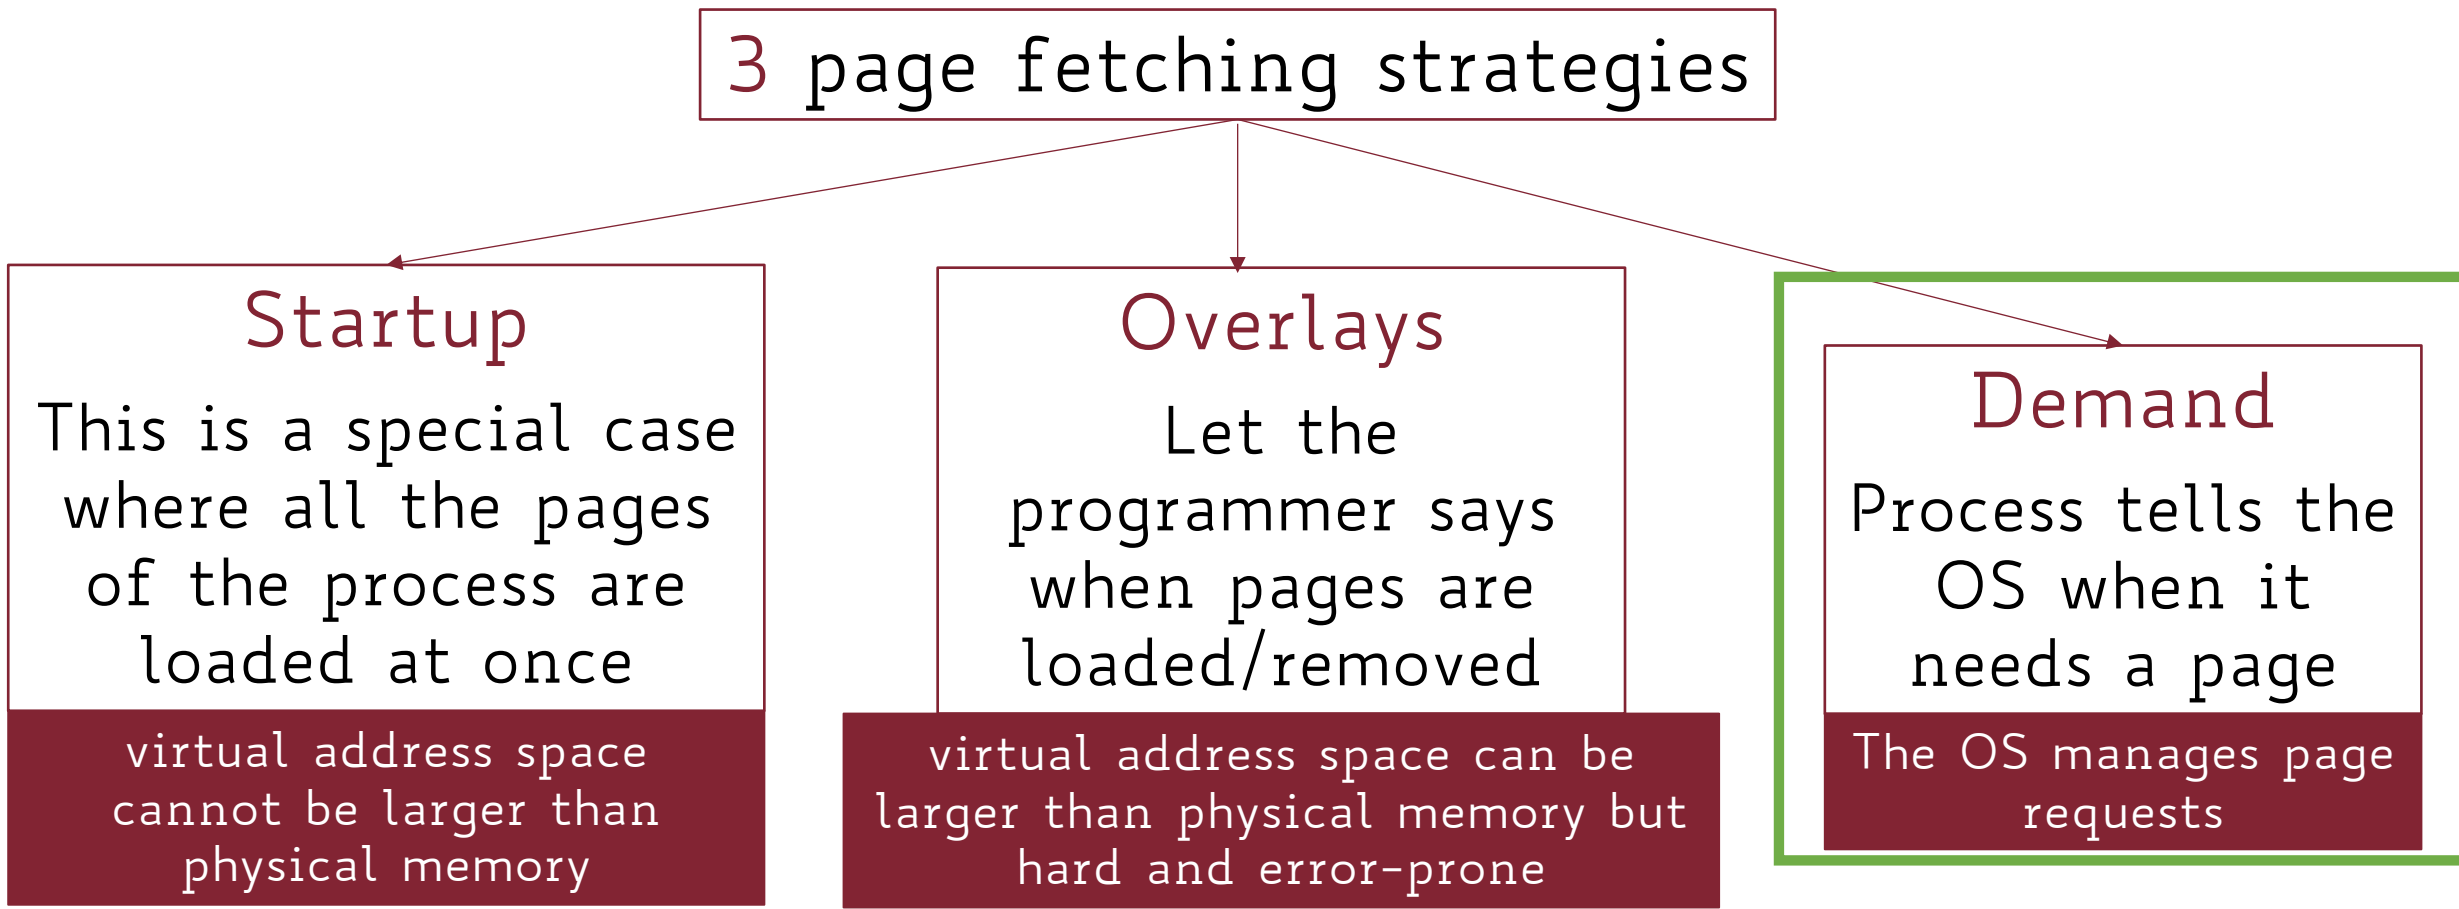
## Demand

Process tells the OS when it needs a page

The OS manages page requests

# Page Fetching Strategies

3 page fetching strategies



```
graph TD; A[3 page fetching strategies] --> B[Startup]; A --> C[Overlays]; A --> D[Demand];
```

## Startup

This is a special case where all the pages of the process are loaded at once

virtual address space cannot be larger than physical memory

## Overlays

Let the programmer say when pages are loaded/removed

virtual address space can be larger than physical memory but hard and error-prone

## Demand

Process tells the OS when it needs a page

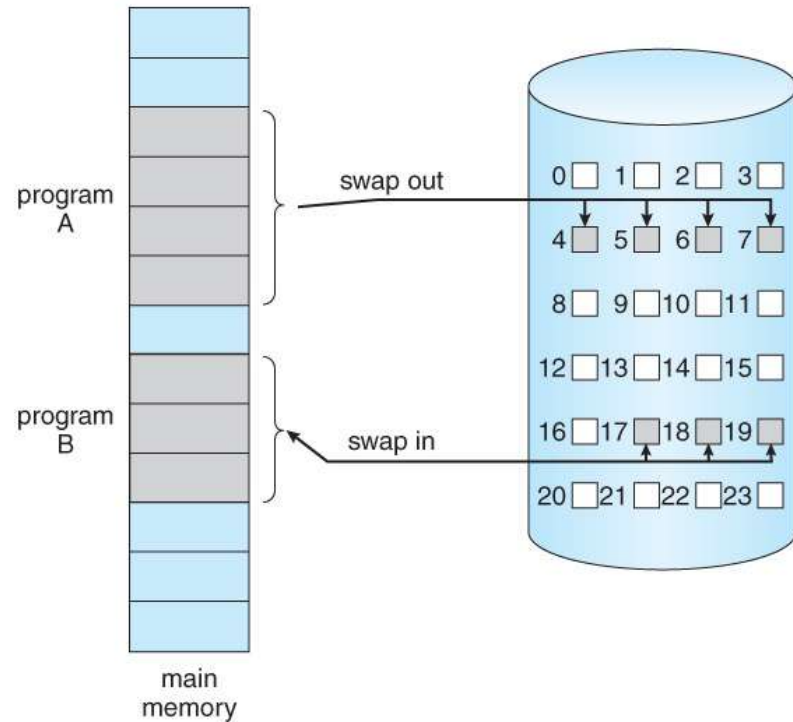
The OS manages page requests

Most modern OSs use demand fetching

# (Pure) Demand Paging


- When a process starts up, **none** of its pages are loaded
- Rather, a page is swapped in only when the process references it (upon a page fault)
- This is termed a **lazy swapper** or **pager**
- Opposite of loading all the pages at process startup!

# Prefetching



The pager guesses when pages will be needed and load them ahead of time


# Prefetching



The pager guesses when pages will be needed and load them ahead of time

Trying to avoid page faults

# Prefetching




The pager guesses when pages will be needed and load them ahead of time

Trying to avoid page faults

Requires predicting the future → very hard!

# Prefetching



The pager guesses when pages will be needed and load them ahead of time

Trying to avoid page faults

Requires predicting the future → very hard!

Possible approach: upon page fault, load many pages instead of only the faulty one  
works if program accesses memory sequentially



# Swap Space

- A portion of the disk reserved for storing pages that are evicted from memory

# Swap Space

- A portion of the disk reserved for storing pages that are evicted from memory
- May be not part of the actual disk file system

# Swap Space

- A portion of the disk reserved for storing pages that are evicted from memory
- May be not part of the actual disk file system
- On Linux there exists a dedicated swap partition (on disk)
  - no actual files are stored in that partition

# Swap Space

- A portion of the disk reserved for storing pages that are evicted from memory
- May be not part of the actual disk file system
- On Linux there exists a dedicated swap partition (on disk)
  - no actual files are stored in that partition
- On Mac, instead, swap space is part of the file system (swapfiles)

# Swap Out

- When a page needs to be swapped out, it will be generally copied to disk

# Swap Out

- When a page needs to be swapped out, it will be generally copied to disk
- The pages for a process are divided into **2 groups**:
  - **Code** (read-only)
  - **Data** (initialized/uninitialized)

# Swap Out

- When a page needs to be swapped out, it will be generally copied to disk
- The pages for a process are divided into **2 groups**:
  - **Code** (read-only)
  - **Data** (initialized/uninitialized)
- Depending on which kind of page is removed, different optimizations may apply upon page swap-out

# Swap Out Optimizations

- **Code** page (read-only):
  - Code content does not change!
  - Just remove and load it back from executable file stored on disk
  - Make use of the filesystem



# Swap Out Optimizations

- **Code** page (read-only):
  - Code content does not change!
  - Just remove and load it back from executable file stored on disk
  - Make use of the filesystem
- **Data** page:
  - Data content does actually change!
  - Save it to a separate paging file, so that no changes are lost when it will be loaded in the future
  - Need to use the dedicated swap space

# Page Replacement: Motivation

- On a page fault, we need to load a page from disk into memory

# Page Replacement: Motivation

- On a page fault, we need to load a page from disk into memory
- If physical memory has still free frames, the page can be safely loaded into one of those

# Page Replacement: Motivation

- On a page fault, we need to load a page from disk into memory
- If physical memory has still free frames, the page can be safely loaded into one of those
- If physical memory is full, a frame must be swapped out to make room for the swap-in page

# Page Replacement: Motivation

- On a page fault, we need to load a page from disk into memory
- If physical memory has still free frames, the page can be safely loaded into one of those
- If physical memory is full, a frame must be swapped out to make room for the swap-in page
- Several algorithms to select the page to evict from memory

# Page Replacement Algorithms

- **Random:** pick any page at random (works surprisingly well!)

# Page Replacement Algorithms

- **Random:** pick any page at random (works surprisingly well!)
- **FIFO (First-In-First-Out):** throw out the page that has been in memory for longest time (i.e., the oldest)
  - Easy to implement but may remove frequently accessed pages

# Page Replacement Algorithms

- **Random:** pick any page at random (works surprisingly well!)
- **FIFO (First-In-First-Out):** throw out the page that has been in memory for longest time (i.e., the oldest)
  - Easy to implement but may remove frequently accessed pages
- **MIN (OPT):** remove the page that will not be accessed for the longest time (provably optimal [Belady 1966])
  - Needs to predict the future → very hard!



# Page Replacement Algorithms

- **Random**: pick any page at random (works surprisingly well!)
- **FIFO (First-In-First-Out)**: throw out the page that has been in memory for longest time (i.e., the oldest)
  - Easy to implement but may remove frequently accessed pages
- **MIN (OPT)**: remove the page that will not be accessed for the longest time (provably optimal [Belady 1966])
  - Needs to predict the future → very hard!
- **LRU (Least Recently Used)**: approximation of MIN, remove the page that has not been used in the longest time
  - Assumes the past is a good predictor of the future (not always true!)

# FIFO Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$											
$F_2$											
$F_3$											

How many page faults (denoted by \*)?

# FIFO Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$											
$F_2$											
$F_3$											

Initially, no frame is loaded in memory at all  
(pure demand paging)

# FIFO Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$											
$F_2$											
$F_3$											

Virtual address within page A is referenced

# FIFO Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$											
$F_2$											
$F_3$											

Virtual address within page A is referenced

page fault

# FIFO Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*										
$F_2$											
$F_3$											

Virtual address within page A is referenced → page fault → A loaded

FIFO = A

# FIFO Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A									
$F_2$											
$F_3$											

Virtual address within page B is referenced

FIFO = A

# FIFO Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A									
$F_2$											
$F_3$											

Virtual address within page B is referenced

page fault

FIFO = A



# FIFO Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A									
$F_2$		B*									
$F_3$											

Virtual address within page B is referenced

page fault



B loaded

# FIFO Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A								
$F_2$		B*	B								
$F_3$											

Virtual address within page C is referenced

FIFO = A → B

# FIFO Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A								
$F_2$		B*	B								
$F_3$											

Virtual address within page C is referenced

page fault

FIFO = A → B

# FIFO Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A								
$F_2$		B*	B								
$F_3$			C*								

Virtual address within page C is referenced

page fault



C loaded

FIFO = A → B → C

# FIFO Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A							
$F_2$		B*	B	B							
$F_3$			C*	C							

Virtual address within page A is referenced

A is already loaded

FIFO = A → B → C

# FIFO Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A						
$F_2$		B*	B	B	B						
$F_3$			C*	C	C						

Virtual address within page B is referenced

B is already loaded

FIFO = A → B → C

# FIFO Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	A					
$F_2$		B*	B	B	B	B					
$F_3$			C*	C	C	C					

Virtual address within page D is referenced

FIFO = A → B → C

# FIFO Page Replacement: Example

3 physical frames:  $F_1, F_2, F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	A					
$F_2$		B*	B	B	B	B					
$F_3$			C*	C	C	C					

Virtual address within page D is referenced

page fault

FIFO = A → B → C



# FIFO Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	D*					
$F_2$		B*	B	B	B	B					
$F_3$			C*	C	C	C					

Virtual address within page D is referenced

page fault

A replaced  
D loaded

FIFO = B → C → D

# FIFO Page Replacement: Example

3 physical frames:  $F_1, F_2, F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	D*	D				
$F_2$		B*	B	B	B	B	B				
$F_3$			C*	C	C	C	C				

Virtual address within page A is referenced

FIFO = B → C → D

# FIFO Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	D*	D				
$F_2$		B*	B	B	B	B	B				
$F_3$			C*	C	C	C	C				

Virtual address within page A is referenced

page fault

FIFO = B → C → D

# FIFO Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	D*	D				
$F_2$		B*	B	B	B	B	A*				
$F_3$			C*	C	C	C	C				

Virtual address within page A is referenced

page fault

B replaced  
A loaded

FIFO = C → D → A

# FIFO Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	D*	D	D			
$F_2$		B*	B	B	B	B	A*	A			
$F_3$			C*	C	C	C	C	C			

Virtual address within page D is referenced

D is already loaded

FIFO = C → D → A

# FIFO Page Replacement: Example

3 physical frames:  $F_1, F_2, F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	D*	D	D	D	C*	C
$F_2$		B*	B	B	B	B	A*	A	A	A	A
$F_3$			C*	C	C	C	C	C	B*	B	B

Eventually, we get a total of 7 page faults

# MIN Page Replacement: Example

3 physical frames:  $F_1, F_2, F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$											
$F_2$											
$F_3$											

How many page faults (denoted by \*)?

# MIN Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$											
$F_2$											
$F_3$											

Initially, no frame is loaded in memory at all  
(pure demand paging)



# MIN Page Replacement: Example

3 physical frames:  $F_1, F_2, F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A						
$F_2$		B*	B	B	B						
$F_3$			C*	C	C						

Up to this point, the same as FIFO

# MIN Page Replacement: Example

3 physical frames:  $F_1, F_2, F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	A					
$F_2$		B*	B	B	B	B					
$F_3$			C*	C	C	C					

Virtual address within page D is referenced

# MIN Page Replacement: Example

3 physical frames:  $F_1, F_2, F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	A					
$F_2$		B*	B	B	B	B					
$F_3$			C*	C	C	C					

Virtual address within page D is referenced

page fault

# MIN Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	A					
$F_2$		B*	B	B	B	B					
$F_3$			C*	C	C	C					

Virtual address within page D is referenced

**page fault**

What's the page that will be requested the furthest away?

# MIN Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	A					
$F_2$		B*	B	B	B	B					
$F_3$			C*	C	C	D*					

Virtual address within page D is referenced → **page fault** → C replaced  
D loaded

# MIN Page Replacement: Example

3 physical frames:  $F_1, F_2, F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	A	A	A	A		
$F_2$		B*	B	B	B	B	B	B	B		
$F_3$			C*	C	C	D*	D	D	D		

Up to this point, no more page faults

# MIN Page Replacement: Example

3 physical frames:  $F_1, F_2, F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	A	A	A	A	A	
$F_2$		B*	B	B	B	B	B	B	B	B	
$F_3$			C*	C	C	D*	D	D	D	D	

Virtual address within page C is referenced

# MIN Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	A	A	A	A	A	
$F_2$		B*	B	B	B	B	B	B	B	B	
$F_3$			C*	C	C	D*	D	D	D	D	

Virtual address within page C is referenced

page fault

What's the page that will be requested the furthest away?



# MIN Page Replacement: Example

3 physical frames:  $F_1, F_2, F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	A	A	A	A	A	
$F_2$		B*	B	B	B	B	B	B	B	C*	
$F_3$			C*	C	C	D*	D	D	D	D	

Virtual address within page C is referenced

page fault

B replaced  
C loaded

B or D will be requested the furthest away (surely not A):  
pick one (e.g., B)

# MIN Page Replacement: Example

3 physical frames:  $F_1, F_2, F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	A	A	A	A	A	A
$F_2$		B*	B	B	B	B	B	B	B	C*	C
$F_3$			C*	C	C	D*	D	D	D	D	D

Eventually, we get a total of 5 page faults

# LRU Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$											
$F_2$											
$F_3$											

How many page faults (denoted by \*)?

# LRU Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$											
$F_2$											
$F_3$											

Initially, no frame is loaded in memory at all  
(pure demand paging)

# LRU Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A						
$F_2$		B*	B	B	B						
$F_3$			C*	C	C						

Up to this point, the same as FIFO

# LRU Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	A					
$F_2$		B*	B	B	B	B					
$F_3$			C*	C	C	C					

Virtual address within page D is referenced

page fault

# LRU Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	A					
$F_2$		B*	B	B	B	B					
$F_3$			C*	C	C	C					

Virtual address within page D is referenced

page fault

We can't look forward anymore!

# LRU Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	A					
$F_2$		B*	B	B	B	B					
$F_3$			C*	C	C	D*					

Virtual address within page D is referenced

page fault

C replaced  
D loaded

C is the page that has not been used for the longest time in the past



# LRU Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	A	A	A	A		
$F_2$		B*	B	B	B	B	B	B	B		
$F_3$			C*	C	C	D*	D	D	D		

Up to this point, no more page faults

# LRU Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	A	A	A	A	A	
$F_2$		B*	B	B	B	B	B	B	B	B	
$F_3$			C*	C	C	D*	D	D	D	D	

Virtual address within page C is referenced

# LRU Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	A	A	A	A	A	
$F_2$		B*	B	B	B	B	B	B	B	B	
$F_3$			C*	C	C	D*	D	D	D	D	

Virtual address within page C is referenced

page fault

We can't look forward anymore!

# LRU Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	A	A	A	A	C*	
$F_2$		B*	B	B	B	B	B	B	B	B	
$F_3$			C*	C	C	D*	D	D	D	D	

Virtual address within page C is referenced

page fault

A replaced  
C loaded

A is the page that has not been used for the longest time in the past

# LRU Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	A	A	A	A	C*	C
$F_2$		B*	B	B	B	B	B	B	B	B	B
$F_3$			C*	C	C	D*	D	D	D	D	D

Virtual address within page A is referenced

# LRU Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	A	A	A	A	C*	C
$F_2$		B*	B	B	B	B	B	B	B	B	B
$F_3$			C*	C	C	D*	D	D	D	D	D

Virtual address within page A is referenced

page fault

We can't look forward anymore!

# LRU Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	A	A	A	A	C*	C
$F_2$		B*	B	B	B	B	B	B	B	B	B
$F_3$			C*	C	C	D*	D	D	D	D	A*

Virtual address within page A is referenced

page fault

D replaced  
A loaded

D is the page that has not been used for the longest time in the past

# LRU Page Replacement: Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, A, B, D, A, D, B, C, A

	A	B	C	A	B	D	A	D	B	C	A
$F_1$	A*	A	A	A	A	A	A	A	A	C*	C
$F_2$		B*	B	B	B	B	B	B	B	B	B
$F_3$			C*	C	C	D*	D	D	D	D	A*

Eventually, we get a total of 6 page faults



# LRU Page Replacement: (An Unlucky) Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, D, A, B, C, D, A, B, C

	A	B	C	D	A	B	C	D	A	B	C
$F_1$											
$F_2$											
$F_3$											

How many page faults (denoted by \*)?

# LRU Page Replacement: (An Unlucky) Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, D, A, B, C, D, A, B, C

	A	B	C	D	A	B	C	D	A	B	C
$F_1$	A*	A	A								
$F_2$		B*	B								
$F_3$			C*								

# LRU Page Replacement: (An Unlucky) Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, D, A, B, C, D, A, B, C

	A	B	C	D	A	B	C	D	A	B	C
$F_1$	A*	A	A	D*							
$F_2$		B*	B	B							
$F_3$			C*	C							

# LRU Page Replacement: (An Unlucky) Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, D, A, B, C, D, A, B, C

	A	B	C	D	A	B	C	D	A	B	C
$F_1$	A*	A	A	D*	D						
$F_2$		B*	B	B	A*						
$F_3$			C*	C	C						

# LRU Page Replacement: (An Unlucky) Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, D, A, B, C, D, A, B, C

	A	B	C	D	A	B	C	D	A	B	C
$F_1$	A*	A	A	D*	D	D					
$F_2$		B*	B	B	A*	A					
$F_3$			C*	C	C	B*					

# LRU Page Replacement: (An Unlucky) Example

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, D, A, B, C, D, A, B, C

	A	B	C	D	A	B	C	D	A	B	C
$F_1$	A*	A	A	D*	D	D	C*				
$F_2$		B*	B	B	A*	A	A				
$F_3$			C*	C	C	B*	B				

# LRU Page Replacement: (An Unlucky) Example

3 physical frames:  $F_1, F_2, F_3$

4 virtual pages: A, B, C, D

Reference sequence of pages: A, B, C, D, A, B, C, D, A, B, C

	A	B	C	D	A	B	C	D	A	B	C
$F_1$	A*	A	A	D*	D	D	C*	C	C	B*	B
$F_2$		B*	B	B	A*	A	A	D*	D	D	C*
$F_3$			C*	C	C	B*	B	B	A*	A	A

Eventually, we get a total of 11 page faults

# Page Replacement: What If We Add Memory?

- Does adding memory always reduce the number of page faults?
- Intuitively, it would seem so...
- The answer, in fact, depends on the page replacement algorithm
- Let's see this with an example, using FIFO page replacement



# FIFO Page Replacement: Example

5 virtual pages: A, B, C, D, E

3 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$

Scenario 1

4 physical frames:  $F_1$ ,  $F_2$ ,  $F_3$ ,  $F_4$

Scenario 2

Reference sequence of pages: A, B, C, D, A, B, E, A, B, C, D, E

# FIFO Page Replacement: Example

	A	B	C	D	A	B	E	A	B	C	D	E
F <sub>1</sub>	A*	A	A	D*	D	D	E*	E	E	E	E	E
F <sub>2</sub>		B*	B	B	A*	A	A	A	A	C*	C	C
F <sub>3</sub>			C*	C	C	B*	B	B	B	B	D*	D
F <sub>1</sub>	A*	A	A	A	A	A	E*	E	E	E	D*	D
F <sub>2</sub>		B*	B	B	B	B	B	A*	A	A	A	E*
F <sub>3</sub>			C*	C	C	C	C	C	B*	B	B	B
F <sub>4</sub>				D*	D	D	D	D	D	C*	C	C

# FIFO Page Replacement: Example

	A	B	C	D	A	B	E	A	B	C	D	E
F <sub>1</sub>	A*	A	A	D*	D	D	E*	E	E	E	E	E
F <sub>2</sub>		B*	B	B	A*	A	A	A	A	C*	C	C
F <sub>3</sub>			C*	C	C	B*	B	B	B	B	D*	D
F <sub>1</sub>	A*	A	A	A	A	A	E*	E	E	E	D*	D
F <sub>2</sub>		B*	B	B	B	B	B	A*	A	A	A	E*
F <sub>3</sub>			C*	C	C	C	C	C	B*	B	B	B
F <sub>4</sub>				D*	D	D	D	D	D	C*	C	C

10 page faults

11 page faults

## Belady's Anomaly

Adding page frames may cause more page faults with some algorithms

# LRU Page Replacement: Example

	A	B	C	D	A	B	E	A	B	C	D	E
F <sub>1</sub>	A*	A	A	D*	D	D	E*	E	E	C*	C	C
F <sub>2</sub>		B*	B	B	A*	A	A	A	A	A	D*	D
F <sub>3</sub>			C*	C	C	B*	B	B	B	B	B	B
F <sub>1</sub>	A*	A	A	A	A	A	A	A	A	A	A	E*
F <sub>2</sub>		B*	B	B	B	B	B	B	B	B	B	B
F <sub>3</sub>			C*	C	C	C	E*	E	E	E	D*	D
F <sub>4</sub>				D*	D	D	D	D	D	C*	C	C

9 page faults

8 page faults

With LRU, adding page frames **always** decreases the number of page faults

# LRU Page Replacement: Example

	A	B	C	D	A	B	E	A	B	C	D	E
F <sub>1</sub>	A*	A	A	D*	D	D	E*	E	E	C*	C	C
F <sub>2</sub>		B*	B	B	A*	A	A	A	A	A	D*	D
F <sub>3</sub>			C*	C	C	B*	B	B	B	B	B	B
F <sub>1</sub>	A*	A	A	A	A	A	A	A	A	A	A	E*
F <sub>2</sub>		B*	B	B	B	B	B	B	B	B	B	B
F <sub>3</sub>			C*	C	C	C	E*	E	E	E	D*	D
F <sub>4</sub>				D*	D	D	D	D	D	C*	C	C

9 page faults

8 page faults

With LRU, adding page frames *always* decreases the number of page faults

Why?

# LRU Page Replacement: Example

	A	B	C	D	A	B	E	A	B	C	D	E
F <sub>1</sub>	A*	A	A	D*	D	D	E*	E	E	C*	C	C
F <sub>2</sub>		B*	B	B	A*	A	A	A	A	A	D*	D
F <sub>3</sub>			C*	C	C	B*	B	B	B	B	B	B
F <sub>1</sub>	A*	A	A	A	A	A	A	A	A	A	A	E*
F <sub>2</sub>		B*	B	B	B	B	B	B	B	B	B	B
F <sub>3</sub>			C*	C	C	C	E*	E	E	E	D*	D
F <sub>4</sub>				D*	D	D	D	D	D	C*	C	C

At each point in time 4-frame memory contains a subset of 3-frame

Can't do any worst!

# Page Replacement: Summary

- **FIFO** is easy to implement but may lead to too many page faults
- May suffer from Belady's Anomaly

# Page Replacement: Summary

- **MIN** is the optimal choice but cannot be used in practice since future memory references are never known in advance



# Page Replacement: Summary

- LRU is a fair approximation of MIN assuming the past is a good predictor of the future
  - Exploits the locality reference (small working set that fits in memory)
  - Works poorly when the locality reference doesn't hold (large working set)